

Chapter 41. Naive Bayes Classification

Bayesian classification methods rely on Bayes' Theorem. Bayes' Theorem describe the relationship of conditional probabilities in light of prior information.

If we are interested in finding the probability that an event L occurs given that we have observe certain *features* in the data, we can express it as $P(L|features)$. Bayes' Theorem tells us this can be calculated as:

$$P(L|features) = \frac{P(features|L)P(L)}{P(features)}$$

If our goal is to decide between two labels (binary classification) L_1 and L_2 , then we can compute it by the ratio of posterior probabilities:

$$\frac{P(L_1|features)}{P(L_2|features)} = \frac{P(features|L_1)P(L_1)}{P(features|L_2)P(L_2)}$$

All we need is a model by which we can compute $P(features|L_i)$ for each label. Such model is called *generative model*. It hypothesises the random process that generates the data. Specifying this model is the main piece of training a Bayesian classifier.

This is where the name "naive" comes from: making naive assumptions about the generative process of the data to find an approximation to it.

In [1]:

```
%matplotlib inline
# basic imports
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('whitegrid')
```


Gaussian Naive Bayes

Gaussian naive Bayes assumes the data is drawn from a *Gaussian distribution* (aka Normal distribution).

Let us generate some data:

In [3]:

```
from sklearn.datasets import make_blobs  
  
X, y = make_blobs(100, 2, centers=2, random_state=2, cluster_std=1.5)  
plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='RdBu');
```



Fitting a GaussianNB model in sklearn to find the generative of each class:

In [4]:

```
from sklearn.naive_bayes import GaussianNB
model = GaussianNB()
model.fit(X,y);
```

Generating new data to predict labels:


In [6]:

```
rng = np.random.RandomState(0)
Xnew = [-6, -14] + [14, 18] * rng.rand(2000, 2)
ynew = model.predict(Xnew)
```

Let's plot this new data to visualize the **decision boundary** learned from the data:

In [7]:

```
plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='RdBu')
lim = plt.axis()
plt.scatter(Xnew[:, 0], Xnew[:, 1], c=ynew, s=20, cmap='RdBu', alpha=0.1)
plt.axis(lim);
```



The Bayesian formalism allows for probabilistic classification, which we can obtain using `predict_proba`:

In [8]:

```
yprob = model.predict_proba(Xnew)  
yprob[-8:].round(2)
```

Out[8]:

```
array([[0.89, 0.11],
       [1.   , 0.   ],
       [1.   , 0.   ],
       [1.   , 0.   ],
       [1.   , 0.   ],
       [1.   , 0.   ],
       [0.   , 1.   ],
       [0.15, 0.85]])
```

The above result is the probability that each point belongs in classes 1 and 2, respectively. For instance, the first data point was predicted to have a probability of 0.89 (89%) of being in class 1 versus 0.11 (11%) of being in class 2.

Notice that `GaussianNB` is non-parametric. We do not learn any parameters from training, we simply make an assumption about the data distribution and fit the data according to this assumption.

Chapter 44: Decision Trees and Random Forests


Random Forests is another nonparametric algorithm built of multiple *Decision Trees*. Random Forests is called an **ensemble method** because it relies on multiple, smaller models *voting* on an output. Votes are aggregated by a majority count, which is taken as the output of a Random Forest model.

The sum can be greater than the parts: Often, the majority voting of multiple decision trees in a random forests model can outperform the vote of a single decision tree.

Decision Trees

Decision trees are intuitive ways to classify an object: ask a series of questions about it to close in on the classification.

For example, a decision tree to classify animals:



Creating a Decision Tree


Generating synthetic data:

In [9]:


```
from sklearn.datasets import make_blobs

X, y = make_blobs(n_samples=300, centers=4,
                  random_state=0, cluster_std=1.0)

plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='rainbow');
```



A decision tree will iteratively split the data along one or the other axis according to some criterion:



A region that contains only points of a single class is not split again.

Fitting a DecisionTreeClassifier:

In [10]:

```
from sklearn.tree import DecisionTreeClassifier
tree = DecisionTreeClassifier().fit(X, y)
```

The following utility help us visualize a decision tree classifier:

In [11]:

```
def visualize_classifier(model, X, y, ax=None, cmap='rainbow'):
    ax = ax or plt.gca()

    # Plot the training points
    ax.scatter(X[:, 0], X[:, 1], c=y, s=30, cmap=cmap,
               clim=(y.min(), y.max()), zorder=3)
    ax.axis('tight')
    ax.axis('off')
    xlim = ax.get_xlim()
    ylim = ax.get_ylim()


    # fit the estimator
    model.fit(X, y)
    xx, yy = np.meshgrid(np.linspace(*xlim, num=200),
                          np.linspace(*ylim, num=200))
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)

    # Create a color plot with the results
    n_classes = len(np.unique(y))
```

```
contours = ax.contourf(xx, yy, Z, alpha=0.3,  
                      levels=np.arange(n_classes + 1) - 0.5,  
                      cmap=cmap, zorder=1)  
  
ax.set(xlim=xlim, ylim=ylim)
```

In [12]:

```
visualize_classifier(DecisionTreeClassifier(), X, y)
```




Notice how decision trees may generate strangely shaped decision boundaries. This occurs because of *overfitting* and can be controlled by the `max_depth` hyperparameter of `DecisionTreeClassifier`.

Decision Trees and Overfitting

It is very easy for the decision tree to go deep and overfit to the data.

Let us look at two trees trained on different subsets of the data:




One observation is that combining information from both trees might lead to better results.

Ensemble of Estimators: Random Forests

The notion that multiple overfitting estimators can be combined to reduce the effect of overfitting is the underlying principle of *Random Forests*: An ensemble of decision trees.

In [16]:

```
from sklearn.ensemble import RandomForestClassifier  
  
model = RandomForestClassifier(n_estimators=100, random_state=0)  
visualize_classifier(model, X, y);
```



Example: Random Forests for Classifying Digits

Let's use Random Forests on the handwritten digits dataset:

In [13]:

```
from sklearn.datasets import load_digits
digits = load_digits()
digits.keys()
```

Out[13]:

```
dict_keys(['data', 'target', 'frame', 'feature_names', 'target_names', 'images', 'DESCR'])
```


Let us visualize the digits:

In [14]:

```
# set up the figure
fig = plt.figure(figsize=(6, 6)) # figure size in inches
fig.subplots_adjust(left=0, right=1, bottom=0, top=1,
                    hspace=0.05, wspace=0.05)

# plot the digits: each image is 8x8 pixels
for i in range(64):
    ax = fig.add_subplot(8, 8, i + 1, xticks=[], yticks=[])
    ax.imshow(digits.images[i], cmap=plt.cm.binary, interpolation='nearest')

    # label the image with the target value
    ax.text(0, 7, str(digits.target[i]))
```



Setting up a Random Forests classifier:

In [17]:

```

from sklearn.model_selection import train_test_split

Xtrain, Xtest, ytrain, ytest = train_test_split(digits.data, digits.target,
                                                random_state=0)

model = RandomForestClassifier(n_estimators=1000)
model.fit(Xtrain, ytrain)
ypred = model.predict(Xtest)

```

In []:

```

from sklearn import metrics
print(metrics.classification_report(ypred, ytest))

```

	precision	recall	f1-score	support
0	1.00	0.97	0.99	38
1	0.98	0.95	0.97	44
2	0.95	1.00	0.98	42
3	0.98	0.96	0.97	46
4	0.97	1.00	0.99	37
5	0.98	0.96	0.97	49
6	1.00	1.00	1.00	52
7	1.00	0.96	0.98	50
8	0.94	0.98	0.96	46
9	0.96	0.98	0.97	46
accuracy			0.98	450
macro avg	0.98	0.98	0.98	450
weighted avg	0.98	0.98	0.98	450

How well did the model do? We will use Scikit-Learn's `classification_report`, which outputs the *precision*, *recall*, *f1-score*, *support* for each class (digit, in this case). It also reports the overall *accuracy*, *macro average* (unweighted) and *weighted average*.

Metric	Description	Definition
Precision	The ratio of positive predictions to all predictions	$\begin{aligned} & \textit{Precision} \\ &= \frac{TP}{TP+FP} \end{aligned}$
Recall	The ratio of positive predictions to all positive samples	$\begin{aligned} & \textit{Recall} \\ &= \frac{TP}{TP+FN} \end{aligned}$
F1-score	The harmonic mean of <i>precision</i> and <i>recall</i>	$\begin{aligned} & F_1 \\ &= 2 \frac{\textit{Precision} \times \textit{Recall}}{\textit{Precision} + \textit{Recall}} \end{aligned}$

Metric	Description	Definition
Accuracy	The ratio of samples predicted correctly	$\text{Accuracy} = \frac{\sum_i^n \hat{y}_i = y_i}{n}$
Support	The number of data in each class	

Where n is the number of samples, $\hat{y}_i$ is the predicted class, y_i is the true class, TP is the number of true-positive outcomes, FP are false-positive outcomes and FN are false-negative outcomes.

Why Precision, Recall and F_1 ?

Suppose you are developing a model to predict if a patient has a very rare disease that only occurs with a frequency of 1 to 100,000. A uniform sample of the population will, thus, have 1 patient with the disease for every 100,000 individuals.

If you train a model on such dataset, it may achieve very high accuracy by simply predicting that no one has the disease, resulting in an accuracy of 0.999. However, if we measure precision and

recall, we would get $Precision = \frac{0}{99,999}$ and $Recall = \frac{0}{1}$, resulting in $F_1 = 0$.

F_1 is a stronger indicator of the model's performance because it accounts for class imbalance by combining precision and recall.

Exercise

What would be the precision and recall in the above if the model did the opposite and predicted every patient had the disease?

Confusion matrix

A confusion matrix shows how well a model classifies data from each labels, and the errors it makes:

In []:

```
from sklearn.metrics import confusion_matrix
import seaborn as sns
mat = confusion_matrix(ytest, ypred)
sns.heatmap(mat.T, square=True, annot=True, fmt='d',
            cbar=False, cmap='Blues')
plt.xlabel('true label')
plt.ylabel('predicted label');
```

0	37	0	1	0	0	0	0	0	0	0
1	0	42	1	0	0	0	0	0	1	0
2	0	0	42	0	0	0	0	0	0	0
3	0	0	0	44	0	0	0	0	1	1
4	0	0	0	0	37	0	0	0	0	0
5	0	1	0	0	0	47	0	0	0	1
6	0	0	0	0	0	0	52	0	0	0
7	0	0	0	0	1	0	0	48	1	0
8	0	0	0	1	0	0	0	0	45	0
9	0	0	0	0	0	1	0	0	0	45
	0	1	2	3	4	5	6	7	8	9

predicted label

true label


Chapter 43: Support Vector Machines

A Support Vector Machine (SVM) is a type of discriminative classifier. It predicts classes by learning a function that describes a decision boundary between classes.

Consider the following data of well-separated points:

In [18]:

```
from sklearn.datasets import make_blobs
X, y = make_blobs(n_samples=50, centers=2,
                  random_state=0, cluster_std=0.60)
plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn');
```




A linear discriminative classifier may attempt to draw a straight line separating the two sets of data:

In [19]:

```
xfit = np.linspace(-1, 3.5)
plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')
plt.plot([0.6], [2.1], 'x', color='red', markeredgewidth=2, markersize=10)

for m, b in [(1, 0.65), (0.5, 1.6), (-0.2, 2.9)]:
    plt.plot(xfit, m * xfit + b, '-k')

plt.xlim(-1, 3.5);
```



These are different separators that perfectly discriminate the training data. However, when a new datapoint is added ("X"), each will assign different labels to it. Which line is the best?

SVM: Maximizing the Margin


The intuition behind SVMs is to include a *margin* around the discriminator that separates data from each class. It would like like this:

In [20]:

```
xfit = np.linspace(-1, 3.5)
plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')

for m, b, d in [(1, 0.65, 0.33), (0.5, 1.6, 0.55), (-0.2, 2.9, 0.2)]:
    yfit = m * xfit + b
    plt.plot(xfit, yfit, '-k')
    plt.fill_between(xfit, yfit - d, yfit + d, edgecolor='none',
                     color='lightgray', alpha=0.5)

plt.xlim(-1, 3.5);
```



Fitting an SVM

Scikit-Learn's support vector classifier (SVC) implements an SVM classifier model:

In [21]:

```
from sklearn.svm import SVC # "Support vector classifier"  
model = SVC(kernel='linear', C=1E10)
```

```
model.fit(X, y)
```

Out[21]:

▼ SVC ⓘ ?
► Parameters

The SVC model from above uses a linear kernel and a very large margin threshold C .

Let us add a helper function to visualize the SVM decision boundary:

In [22]:

```
def plot_svc_decision_function(model, ax=None, plot_support=True):  
    """Plot the decision function for a 2D SVC"""  
    if ax is None:  
        ax = plt.gca()  
    xlim = ax.get_xlim()  
    ylim = ax.get_ylim()  
  
    # create grid to evaluate model  
    x = np.linspace(xlim[0], xlim[1], 30)  
    y = np.linspace(ylim[0], ylim[1], 30)  
    Y, X = np.meshgrid(y, x)
```

```

xy = np.vstack([X.ravel(), Y.ravel()]).T
P = model.decision_function(xy).reshape(X.shape)

# plot decision boundary and margins
ax.contour(X, Y, P, colors='k',
           levels=[-1, 0, 1], alpha=0.5,
           linestyles=['--', '-', '--'])

# plot support vectors
if plot_support:
    ax.scatter(model.support_vectors_[0],
               model.support_vectors_[1],
               s=300, linewidth=1, edgecolors='black',
               facecolors='none');
ax.set_xlim(xlim)
ax.set_ylim(ylim)


```

In [23]:

```

plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')
plot_svc_decision_function(model);

```



Notice that the SVM chose the discriminator that maximizes the margin between the two sets of points. The points touching the margin are named *support vectors*.

A key aspect of SVMs is that only support vectors matter. Any point further from the margin that are on the correct side do not modify the fit.

Non-linear SVM kernels


If the relation between the data at hand is not linear, we may have to use a different kernel for the SVC:

In [24]:

```
from sklearn.datasets import make_circles
X, y = make_circles(100, factor=.1, noise=.1)

clf = SVC(kernel='linear').fit(X, y)

plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')
plot_svc_decision_function(clf, plot_support=False);
```



A *kernel transformation* applies a transform function to the data in order to make it "separable" in at higher dimensions. One such kernels is the *radial basis function*(rbf) kernel. We can specify the kernel to be used with the `kernel` parameter:

In [25]:


```
clf = SVC(kernel='rbf', C=1E6)
clf.fit(X, y)
```

Out[25]:

▼ SVC ⓘ ?
► Parameters

In [26]:

```
plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')  
plot_svc_decision_function(clf)  
plt.scatter(clf.support_vectors_[:, 0], clf.support_vectors_[:, 1],  
            s=300, lw=1, facecolors='none');
```




Softening the Margins

As you may have noticed, we have applied an SVM to a perfectly separable set of points. And our enforcement of the margin was very strict: every point should be on or behind the margin. In reality, such separation may not be possible. For instance, the data may look like:

In [27]:

```
X, y = make_blobs(n_samples=100, centers=2,  
                  random_state=0, cluster_std=1.2)  
plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn');
```



To handle this case, the SVM implementation includes a parameter that "softens" the margin, allowing some points to creep into it if that allows for a better fit.

The hardness of the margin is controlled by parameter C . For very large C , the margin is hard and allows no point within it. For smaller C , the margin allows a few points inside.


Let us compare two different values of C :

In [29]:

```
X, y = make_blobs(n_samples=100, centers=2,
                  random_state=0, cluster_std=0.8)

fig, ax = plt.subplots(1, 2, figsize=(16, 6))
fig.subplots_adjust(left=0.0625, right=0.95, wspace=0.1)

for axi, C in zip(ax, [10.0, 0.1]):
    model = SVC(kernel='linear', C=C).fit(X, y)
    axi.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')
    plot_svc_decision_function(model, axi)
    axi.scatter(model.support_vectors_[:, 0],
                model.support_vectors_[:, 1],
                s=300, lw=1, facecolors='none');
    axi.set_title('C = {0:.1f}'.format(C), size=14)
```



On to the lab

Face recognition using SVMs:



Colin Powell



George W Bush



George W Bush



George W Bush



Hugo Chavez



George W Bush



Junichiro Koizumi



George W Bush



Tony Blair



Ariel Sharon



George W Bush



Donald Rumsfeld



George W Bush



George W Bush



George W Bush